

Srinivas Sivaratri

AI Engineer | LLM evaluation, model debugging, developer tools
srinivassivaratri1122@gmail.com | linkedin.com/in/srinivas-sivaratri-3a0a32162 | github.com/sivaratrisrinivas
srinivassivaratri.is-a.dev | Remote, India (IST, UTC+5:30) | Open to full-time and contract

Summary

AI engineer with an MS in Computer Science. Builds LLM systems and the evaluations that check them, and publishes the numbers, including the ones that lose to a baseline. Upstream work reviewed by better-auth and go-ethereum maintainers.

Skills

AI and LLM: LLM evaluation, token-level tracing and replay, grounding checks, RAG, LoRA fine-tuning, diffusion models, speech-to-text and text-to-speech pipelines

Machine Learning: PyTorch, Hugging Face Transformers, PEFT, sentence-transformers

Languages: Python, TypeScript, Go, JavaScript, SQL, C

Engineering: FastAPI, Flask, Next.js, React, Node.js, Bun, PostgreSQL, Redis, SQLite, Docker, GitHub Actions, pytest, Vitest

Experience

Software Engineer, Independent | Remote May 2025 to present

- Traced an email-enumeration leak in better-auth to its Have I Been Pwned plugin and proposed a request-scoped skip flag with a regression test. Maintainers adopted the approach and shipped it (PR #5998).
- Designed "geth db truncate-freezer" for go-ethereum (PR #31351), revised across three core-maintainer review rounds to batch-process instead of holding about 15 GB in memory. The core team later built its own version.

Software Engineer, Accenture | Bangalore, India July 2021 to July 2022

- Wrote a Python ETL script that automated data processing, saved 50 labor hours per month, and cut errors.
- Implemented a GitHub Actions CI/CD pipeline.

Projects

llmfr (LLM Flight Recorder) | Python, PyTorch, Transformers, SQLite github.com/sivaratrisrinivas/llm-flight-recorder

- Built a CLI that records a local LLM generation token by token, replays it, and names the first step where two runs split, tagging every later difference as fallout. The OpenAI adapter refuses to record when logprobs are missing.
- On pinned Qwen2.5-0.5B demos, it labels a seed-only change as a sampling split at token 0 and a temperature-only change as a decoding split at token 8. Compare takes 0.151 s median on CPU.

Fuzz | TypeScript, Bun, FastAPI, Hugging Face, PEFT github.com/sivaratrisrinivas/fuzz

- Built a browser game where waves scramble a private paragraph, the player rewrites to save pieces, and Qwen2.5 rebuilds the rest. The original text never leaves the browser, and the server rejects any request that contains it.
- Fine-tuned a LoRA adapter on Qwen2.5-0.5B with 288 synthetic examples. On the locked eval, word F1 on intact text rose from 0.746 to 0.938 and failed outputs fell from 12/44 to 8/44. Heavily scrambled inputs did not improve.

Sunkelo | Next.js, TypeScript, Postgres, Redis, Sarvam AI github.com/sivaratrisrinivas/sunkelo

- Built a voice-first product review app in 11 Indian languages. It transcribes speech, scrapes public reviews, has an LLM summarize them, then translates and reads the verdict aloud, streamed over SSE with a Redis cache.
- Built a 379-claim grounding eval showing the first word-overlap metric flagged all 13 supported real claims as unsupported. By hand labels, 6 of 19 real claims lack source support.

SokoFlow | Python, PyTorch, Flask github.com/sivaratrisrinivas/sokoflow

- Trained a diffusion model (CNN board encoder, Transformer denoiser) that generates Sokoban move plans. It solved 6.2% (15/240) of scrambled 8x8 boards against 94.2% for plain BFS search. Traced an earlier 20.8% figure to easy boards and published the correction.

Education

MS, Computer Science, University of Alabama at Birmingham, AL, USA Aug 2023 to May 2025

BS, Computer Science, BV Raju Institute of Technology, India Aug 2017 to June 2021

Boot.dev, coursework, HTTP Servers in Go, RAG, CI/CD with GitHub Actions and Docker, Memory Management in C, SQL